

Lab: Build an Automated CD DevOps Pipeline Using Tekton and OpenShift



Estimated time needed: 90 minutes

Welcome to the **Build an Automated CD DevOps Pipeline Using Tekton and OpenShift** hands-on lab. In this lab, you will create a parameterized Tekton pipeline that will allow you to reuse your pipeline-as-code artifacts, build your application, and deploy it to Kubernetes/OpenShift.

If you forget how any of these exercises work, refer to module 3 of the course [Continuous Integration and Continuous Delivery \(CI/CD\)](#) for detailed instructions on how to complete these tasks.

Objectives

In this lab, you will:

- Check out your capstone project code created in earlier modules
- Create a workspace and add a cd-pipeline in the workspace
- Create a lint task using flake8
- Create a test task using the Nose test runner
- Create a build task to build a container image
- Create a deploy task to deploy to an OpenShift cluster
- Apply the changes and run the cd-pipeline

Note: Important Security Information

Welcome to the Cloud IDE with OpenShift. This is where all your development will take place. It has all the tools you will need to use Docker for deploying a PostgreSQL database.

It is important to understand that the lab environment is **ephemeral**. It only lives for a short while and then it will be destroyed. This makes it imperative that you push all changes made to your own GitHub repository so that it can be recreated in a new lab environment any time it is needed.

Also note that this environment is shared and therefore not secure. You should not store any personal information, usernames, passwords, or access tokens in this environment for any purposes.

Your Task

1. If you haven't generated a **GitHub Personal Access Token** you should do so now. You will need it to push code back to your repository. It should have `repo` and `write` permissions, and set to expire in 60 days. When Git prompts you for a password in the Cloud IDE environment, use your Personal Access Token instead.
2. The environment may be recreated at any time so you may find that you have to perform the **Initialize Development Environment** each time the environment is created.

Note on Screenshots

Throughout this lab, you will be prompted to take screenshots and save them on your device. These will be required either to answer graded quiz questions or for submission under **Option 1: AI-Graded Submission and Evaluation** or **Option 2: Peer-Graded Submission and Evaluation** at the end of this course. Your screenshot must have either the `.jpeg` or `.png` extension.

To take screenshots, you can use various free screen-capture tools or your operating system's shortcut keys. For example:

- **Mac:** you can use `Shift + Command + 3` (`⇧ + ⌘ + 3`) on your keyboard to capture your entire screen, or `Shift + Command + 4` (`⇧ + ⌘ + 4`) to capture a window or area. They will be saved as a file on your Desktop.
- **Windows:** you can capture your active window by pressing `Alt + Print Screen` on your keyboard. This command copies an image of your active window to the clipboard. Next, open an image editor, paste the image from your clipboard to the image editor, and save the image.

Initialize Development Environment

Because the Cloud IDE with OpenShift environment is ephemeral, it may be deleted at any time. The next time you come into the lab, a new environment may be created. Unfortunately, this means that you will need to initialize your development environment every time it is recreated. This shouldn't happen too often as the environment can last for several days at a time but when it is removed, this is the procedure to recreate it.

Overview

Each time you need to set up your lab development environment, you will need to run three commands.

Each command will be explained in further detail, one at a time, in the following section.

{your_github_account} represents your GitHub account username.

The commands include:

```
git clone https://github.com/{your_github_account}/devops-capstone-project.git
cd devops-capstone-project
bash ./bin/setup.sh
exit
```

Now, let's discuss each of these commands and explain what needs to be done.

Task Details

Initialize your environment using the following steps:

1. Open a terminal with Terminal -> New Terminal if one isn't open already.
2. Next, use the export GITHUB_ACCOUNT= command to export an environment variable that contains the name of your GitHub account.

Note: Substitute your real GitHub account for the {your_github_account} place holder below:

```
export GITHUB_ACCOUNT={your_github_account}
```

3. Then use the following commands to clone your repository, change into the devops-capstone-project directory, and execute the ./bin/setup.sh command.

```
git clone https://github.com/$GITHUB_ACCOUNT/devops-capstone-project.git
cd devops-capstone-project
bash ./bin/setup.sh
```

You should see the follow at the end of the setup execution:

4. Finally, use the `exit` command to close the current terminal. The environment won't be fully active until you open a new terminal in the next step.

```
exit
```

Validate

In order to validate that your environment is working correctly, you must open a new terminal because the Python virtual environment will only activate when a new terminal is created. You should have ended the previous task by using the `exit` command to exit the terminal.

1. Open a terminal with `Terminal -> New Terminal` and check that everything worked correctly by using the `which python` command:

Your prompt should look like this:



```
(venv) theia:project$
```

Check which Python you are using:

```
which python
```

You should get back:

```
(venv) theia:project$ which python  
/home/theia/venv/bin/python  
(venv) theia:project$ █
```

Check the Python version:

```
python --version
```

You should get back some patch level of Python 3.9:

```
(venv) theia:project$ python --version  
Python 3.9.15 ←  
(venv) theia:project$ █
```

This completes the setup of the development environment. Anytime your environment is recreated, you will need to follow this procedure.

You are now ready to start working.

Exercise 1: Pick Up the Next Story

The first thing you need to do is to go to your kanban board to get a story to work on. Take the first story from the top of the Sprint Backlog, move it to In Progress, assign it to yourself, and read the contents.

Your Task

1. Go to your kanban board and take the next story from the top of the Sprint Backlog. It should be titled: "Create a CD pipeline to automate deployment to Kubernetes".
2. Move the story to In Progress.
3. Open the story and assign it to *yourself*.
4. Read the contents of the story.

Results

The story should look like this:

Create a CD pipeline to automate deployment to Kubernetes

As a developer

I need to create a CD pipeline to automate deployment to Kubernetes
So that developers aren't wasting their time doing it manually

Assumptions

- Use Tekton to define the pipeline
- It should clone, lint, test, build, and deploy the service
- Deployment should be to OpenShift
- It can use a manual trigger for this MVP

Acceptance Criteria

```
Given the CD pipeline has been created
When I trigger the pipeline run
Then I should see the accounts service deployed to OpenShift
```

You are now ready to begin working on your story.

Exercise 2: Overview and Setup

In reading your story, you see that the assumptions state the following:

Assumptions

- Use Tekton to define the pipeline
- It should clone, lint, test, build, and deploy the service
- Deployment should be to OpenShift
- It can use a manual trigger for this MVP

It looks like you will need to create a pipeline task for each of these:

Task Name	Build After	Message
clone	-	Cloning the repo \${params.repo-url}
lint	clone	Calling Flake8 linter
tests	clone	Running unit tests with PyUnit
build	lint, tests	Building image for \${params.repo-url}
deploy	build	Deploying \${params.branch} branch of \${params.repo-url}

Luckily, the developer that worked on this before you started to create a pipeline but didn't get very far. They did, however, create a starter **Pipeline** in `pipeline.yaml` with a `clone` task, and defined a **PersistentVolumeClaim** in the `pvc.yaml` file to define a workspace.

Let's see if what we have already works.

Your Task

1. Change into your project directory: `cd devops-capstone-project`.
2. Create a new branch called `cd-pipeline` to work on in the development environment.
3. Run `nosetests` and make sure that all of the test cases are passing. Fix any that fail before proceeding.
4. Use the `oc create` command on the `tekton/pvc.yaml` file to create a workspace for your pipeline to use.

▼ Click here for the answer.

```
oc create -f tekton/pvc.yaml
```

5. Use the `oc apply` command on the `tekton/tasks.yaml` file to create the tasks for your pipeline to use.

▼ Click here for the answer.

```
oc apply -f tekton/tasks.yaml
```

6. Use the `oc apply` command on the `tekton/pipeline.yaml` file to create the initial pipeline.

▼ Click here for the answer.

```
oc apply -f tekton/pipeline.yaml
```

7. If you edit the `pipeline.yaml` file, you will see that the `clone` task requires the `git-clone` task from Tekton Hub. Use the `tkn hub install task` command to install that now.

▼ Click here for the answer.

```
tkn hub install task git-clone
```

Note: If the above command returns a error due to Tekton Version mismatch, please run the below command to fix this.

▼ Click here for the hint.

```
kubectl apply -f https://raw.githubusercontent.com/tektoncd/catalog/main/task/git-clone/0.9/git-clone.yaml
```

1. Use the `tkn pipeline start` command to start the pipeline and make sure the initial pipeline definitions work.

```
tkn pipeline start cd-pipeline \
-p repo-url="https://github.com/$GITHUB_ACCOUNT/devops-capstone-project.git" \
-p branch="main" \
-w name=pipeline-workspace,claimName=pipelinerun-pvc \
-s pipeline \
--showlog
```

Note: Here `main` is mentioned as the branch parameter since all the code changes are merged to main branch from the previous labs.

Results

You should see the following output that proves that the initial pipeline is working. Now it is safe to start adding more tasks to it.

```
(venv) theia:devops-capstone-project$ tkn pipeline start cd-pipeline \
>     -p repo-url="https://github.com/$GITHUB_ACCOUNT/devops-capstone-project.git" \
>     -p branch="main" \
>     -w name=pipeline-workspace,claimName=piplinerun-pvc \
>     --showlog
PipelineRun started: cd-pipeline-run-ncfzx
Waiting for logs to be available...
[clone : clone] + '[' false '=' true ]
[clone : clone] + '[' false '=' true ]
[clone : clone] + '[' false '=' true ]
[clone : clone] + CHECKOUT_DIR=/workspace/output/
[clone : clone] + '[' true '=' true ]
[clone : clone] + cleandir
[clone : clone] + '[' -d /workspace/output/ ]
[clone : clone] + rm -rf '/workspace/output///*'
[clone : clone] + rm -rf '/workspace/output//..[!.]*'
[clone : clone] + rm -rf '/workspace/output//..?*'
[clone : clone] + test -z
[clone : clone] + test -z
[clone : clone] + test -z
[clone : clone] + git config --global --add safe.directory /workspace/output
[clone : clone] + /ko-app/git-init '-url=https://github.com/rofrano/devops-capstone-project'
fspec=' '-path=/workspace/output/' '-sslVerify=true' '-submodules=true' '-depth=1' '-sparse
[clone : clone] {"level":"info","ts":1666621040.7965918,"caller":"git/git.go:176","msg":"Su
ithub.com/rofrano/devops-capstone-project.git @ 58d1c3182ff9e69e878fd605583e16e8b9506d8a (g
in path /workspace/output/"}
[clone : clone] {"level":"info","ts":1666621040.971369,"caller":"git/git.go:215","msg":"Suc
updated submodules in path /workspace/output/"}
[clone : clone] + cd /workspace/output/
[clone : clone] + git rev-parse HEAD
[clone : clone] + RESULT_SHA=58d1c3182ff9e69e878fd605583e16e8b9506d8a
[clone : clone] + EXIT_CODE=0
[clone : clone] + '[' 0 != 0 ]
```

```
[clone : clone] + l  o  :-  o ]
[clone : clone] + printf '%s' 58d1c3182ff9e69e878fd605583e16e8b9506d8a
[clone : clone] + printf '%s' https://github.com/rofrano/devops-capstone-project.git

(venv) theia:devops-capstone-project$
```

Just to be sure it ran successfully, use `tkn pipelinerun ls` and check the status column, which should show Succeeded.

```
tkn pipelinerun ls
```

It should return:

```
(venv) theia:devops-capstone-project$ tkn pipelinerun ls
NAME          STARTED      DURATION      STATUS
cd-pipeline-run-ncfzx  3 minutes ago  11 seconds  Succeeded
(venv) theia:devops-capstone-project$
```

Exercise 3: Create the Lint Task

The first task that you need to add is the lint task. You are going to use flake8 to lint your code. Luckily, Tekton Hub has a flake8 task that you can install and use:

You can use the following button to open the `pipeline.yaml` file to start editing it:

[Open pipeline.yaml in IDE](#)

Your Task

1. Use the `tkn hub install` command to install the task named `flake8` into your Kubernetes namespace.

You should see output similar to this:

```
(venv) theia:devops-capstone-project$ tkn hub install task flake8
Task flake8(0.1) installed in sn-labs-rofrano namespace
(venv) theia:devops-capstone-project$
```

▼ Click here for the answer.

```
tkn hub install task flake8
```

2. Now you will modify the `pipeline.yaml` file to use the new `flake8` task. Start by copying and pasting the existing `clone` task and change the name to `lint`.

3. In reading the documentation for the `flake8` task, you notice that it requires a workspace named `source`. You should change the workspace name to match. You also need to change the `taskRef` to `flake8`.

▼ Click here for a hint.

```
- name: {name here}
workspaces:
- name: {workspace name here}
  workspace: pipeline-workspace
taskRef:
  name: {task name here}
```

4. You now need to set up the parameters. The `flake8` task takes an optional `image` parameter that you will use, and the `args` parameter for `flake8` to run. Change the existing parameters to use `python:3.9-slim` as the image, and pass the following `args` parameters to `flake8: --count --max-complexity=10 --max-line-length=127 --statistics`.

► Click here for a hint.

5. Finally, you need to define when it should run. Add the `runAfter:` tag and indicate that it should run after the previous `clone` task.

▼ Click here for a hint.

```
- name: {name here}
workspaces:
- name: {workspace name here}
  workspace: pipeline-workspace
taskRef:
  name: {task name here}
params:
- name: image
  value: {image name here}
- name: args
  value: {arguments here}
runAfter:
- {task before}
```

Check Your Answer

▼ Click here for the answer.

```
- name: lint
  workspaces:
    - name: source
      workspace: pipeline-workspace
  taskRef:
    name: flake8
  params:
    - name: image
      value: "python:3.9-slim"
    - name: args
      value: ["--count", "--max-complexity=10", "--max-line-length=127", "--statistics"]
  runAfter:
    - clone
```

Start the Pipeline

Now it's time run the pipeline and make sure that the lint task works.

Apply the changes:

```
oc apply -f tekton/pipeline.yaml
```

Start the pipeline:

```
tkn pipeline start cd-pipeline \
  -p repo-url="https://github.com/$GITHUB_ACCOUNT/devops-capstone-project.git" \
  -p branch="main" \
  -w name=pipeline-workspace,claimName=piplinerun-pvc \
  -s pipeline \
  --showlog
```

Check the Run Status

It may take a long time for the pipeline to run. The logs should display on the screen but if they don't, you can see the pipeline run status by opening a new terminal and listing the pipeline runs with:

```
tkn pipelinerun ls
```

You can check the logs of the last run with:

```
tkn pipelinerun logs --last
```

Commit Your Changes

Because the Cloud IDE is ephemeral, it may timeout and you will lose all of your work. It is strongly recommended that you commit your changes and push them to your remote branch so that you can pull them from GitHub if needed.

```
git commit -am 'added link task'  
git push --set-upstream origin cd-pipeline
```

Now you are ready to move on to the next exercise.

Exercise 4: Create a Test Task

The next stage of the pipeline you need to create is a test task. In your development environment, you use Nose to run your tests using the `nosetests` command. Unfortunately, there are no tasks in the **Tekton Hub** for nosetests, so you will write your own.

In this exercise, you will update the `tasks.yaml` file adding a new task called `nose` that uses the shared workspace for the pipeline and runs nosetests in a `python:3.9-slim` image as a shell script as seen in the CI/CD course video.

Here is a bash script to install the Python requirements and run the nosetests. You can use this as the shell script in your new task:

```
#!/bin/bash
set -e
python -m pip install --upgrade pip wheel
pip install -qr requirements.txt
nosetests -v --with-spec --spec-color
```

You can use the following button to open the `tasks.yaml` file to start editing it:

[Open `tasks.yaml` in IDE](#)

Your Task:

1. First, create a new task in the `tasks.yaml` file. Start by copying and pasting the existing `echo` task and change the name to `nose`.

Remember, each new task must be separated using three dashes — on a separate line.

► [Click here for a hint.](#)

2. Next, it's a good idea to add a description to your task so that others who may use it know what it does. Change the `description:` tag to be: `This task will run nosetests on the provided input.`

► [Click here for a hint.](#)

3. Next, you need to include the workspace that has the code that you want to test. Since `flake8` uses the name `source`, you can use that for consistency. Add a workspace named `source`.

► [Click here for a hint.](#)

4. It might be a good idea to allow the passing in of different arguments to `nosetests`, so create a parameter called `args` just like the `flake8` task has, and give it a `description:`, make the `type:` a `string`, and a `default:` with the verbose flag `-v` as the default.

► [Click here for a hint.](#)

5. Quite often, tests require a test database. It might be a good idea to allow the passing in of a database url. Add a new parameter called `database_uri` just like the `args` parameter. Give it a `description:`, make the `type:` a `string`, and a `default:` with SQLite ("sqlite:///test.db") in case they don't have a separate database server.

► [Click here for a hint.](#)

6. Finally, you need to specify the `steps:`, and there is only one. Give it the name `nosetests`.

► [Click here for a hint.](#)

7. Have it run in a `python:3.9-slim` image.

► [Click here for a hint.](#)

8. Also, specify `workingDir` as the path to the workspace you defined (i.e., `$(workspaces.source.path)`).

► [Click here for a hint.](#)

9. Next, add an environment variable called `DATABASE_URI` that gets its value from the new parameter you just created: `$(params.database_uri)`.

► [Click here for a hint.](#)

10. Then, paste the script from above in the `script:` parameter and substitute `$(params.args)` for the parameters passed into `nosetests` instead of the hard-coded ones.

► [Click here for a hint.](#)

Check Your Answer

Your new task should look similar to this:

- ▶ [Click here for the answer.](#)

Apply Your Changes

Apply these changes to your cluster:

```
oc apply -f tekton/tasks.yaml
```

You should see the following output:

```
(venv) theia:devops-capstone-project$ oc apply -f tekton/tasks.yaml
task.tekton.dev/echo created
task.tekton.dev/nose created
(venv) theia:devops-capstone-project$ █
```

Commit Your Changes

Because the Cloud IDE is ephemeral, it may timeout and you will lose all of your work. It is strongly recommended that you commit your changes and push them to your remote branch so that you can pull them from GitHub if needed.

```
git commit -am 'added nose task'
git push
```

Now you are ready to move on to the next exercise.

Exercise 5: Add the Test Task to the Pipeline

Now it's time to update the pipeline to add the test task that uses the new `nose` task that you have created and then run the pipeline to see it work.

You can use the following button to open the `pipeline.yaml` file to start editing it:

[Open pipeline.yaml in IDE](#)

Your Task

1. Start by editing the `pipeline.yaml` file to add the new test task. Start by copying and pasting the `lint` task and change the name to `tests`.

► [Click here for a hint.](#)

2. Leave the workspace named `source` and change the `taskRef:` from `lint` to reference your new `nose` task.

► [Click here for a hint.](#)

3. Change the `image` parameter to `database_uri` with a value of `sqlite:///test.db`.

► [Click here for a hint.](#)

4. Change the value of the existing `args` parameter to specify the arguments to pass to the tests as `"-v --with-spec --spec-color"`.

► [Click here for a hint.](#)

5. Finally, specify that this task should `runAfter:` the `clone` step.

Notice that you are running `lint` and `tests` in parallel because they do not depend on each other. This will make the pipeline run faster.

► [Click here for a hint.](#)

Check Your Answer

Your new step should look similar to this:

► [Click here for the answer.](#)

Commit Your Changes

Because the Cloud IDE is ephemeral, it may timeout and you will lose all of your work. It is strongly recommended that you commit your changes and push them to your remote branch so that you can pull them from GitHub if needed.

```
git commit -am 'added test pipeline task'  
git push
```

Now you are ready to move on to the next exercise.

Exercise 6: Start the Pipeline

Now it's time run the pipeline and make sure that the `tests` task works. The `tests` task will need a PostgreSQL database to run against. If it isn't running in your OpenShift cluster from the last lab, you must provision it again.

Your Task

1. Use the `oc get svc postgresql` command to see if the PostgreSQL service is defined.

```
oc get svc postgresql
```

If it is not found, use the `oc new-app` command to deploy the `postgresql-ephemeral` template.

```
oc new-app postgresql-ephemeral
```

2. Use `oc get pods` to make sure that the `postgres` service is defined and the `postgres` pod is running.

```
oc get pods
```

You should see the following:

(venv) theia:devops-capstone-project\$ oc get all					
NAME	READY	STATUS	RESTARTS	AGE	
pod/openshift-web-console-8bd9fcf8-2mlw8	2/2	Running	0	2d2h	
pod/openshift-web-console-8bd9fcf8-d8wt5	2/2	Running	0	2d2h	
pod/postgresql-1-deploy	0/1	Completed	0	177m	
pod/postgresql-1-p7rfz	1/1	Running	1 (176m ago)	177m	
NAME	DESIRED	CURRENT	READY	AGE	
replicationcontroller/postgresql-1	1	1	1	177m	
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/openshift-web-console	ClusterIP	172.21.114.10	<none>	8000/TCP	2d2h
service/postgresql	ClusterIP	172.21.42.231	<none>	5432/TCP	177
NAME	READY	UP-TO-DATE	AVAILABLE	AGE	
deployment.apps/openshift-web-console	2/2	2	2	2d2h	
NAME	DESIRED	CURRENT	READY	AGE	
replicaset.apps/openshift-web-console-8bd9fcf8	2	2	2	2d2h	
NAME	REVISION	DESIRED	CURRENT	TRIGGERED	
deploymentconfig.apps.openshift.io/postgresql	1	1	1	config,im	

3. You are now ready to apply the changes to your pipeline.

```
oc apply -f tekton/pipeline.yaml
```

You should see the message: pipeline.tekton.dev/cd-pipeline configured

Start the Pipeline

Make sure that the `GITHUB_ACCOUNT` environment variable is set to the name of your GitHub account, and use the `tkn pipeline start` command to run the pipeline.

```
tkn pipeline start cd-pipeline \
  -p repo-url="https://github.com/$GITHUB_ACCOUNT/devops-capstone-project.git" \
  -p branch="main" \
  -w name=pipeline-workspace,claimName=pipelinerun-pvc \
  -s pipeline \
  --showlog
```

Note that because you specified that `lint` and `tests` can run in parallel, you will see the log output of those two tasks intermixed.

Exercise 7: Create a Build Task

Before you can deploy your application, you need to build a Docker image and push it to an image registry. You search [Tekton Hub](#) for the word "build" and you see there is a task called `buildah` that will build images so you decide to use it in your pipeline to build your code.

Instead of installing it yourself, you first check the ClusterTasks in your cluster to see if it already exists. Luckily, the OpenShift environment you are using already has `buildah` installed as a [**ClusterTask**](#). A ClusterTask is installed cluster-wide by an administrator, and anyone can use it in their pipelines without having to install it themselves.

You can always check what ClusterTasks are installed with this command:

```
tkn clustertask ls
```

Now, you will update the `pipeline.yaml` file to use the new `buildah` task. You can use the following button to open the `pipeline.yaml` file to start editing it:

[Open pipeline.yaml in IDE](#)

Your Task

1. Start by editing the `pipeline.yaml` file to add the new build task. Start by copying and pasting the `tests` task and changing the name to `build`.

► Click here for a hint.

2. The workspace name `source` from the `tests` task is fine, but the `taskRef` needs to be changed to `buildah`, and you must add an additional line below it to indicate that it has a `kind:` of `ClusterTask`.

► Click here for a hint.

3. Delete the `database_uri` parameter. It is not needed for this task.
4. Next, change the `args` parameter to `IMAGE` and set the value to `"$(params.build-image)"` to pick up the `build-image` name from the pipeline parameters.
 - ▶ Click here for a hint.
5. Now, you need to add a parameter named `build-image` to the existing list of parameters at the top of the pipeline under `spec.params`.
 - ▶ Click here for a hint.
6. Finally, you must specify what task or tasks this should run after. Because both `tests` and `lint` ran in parallel, you want to wait until both complete before building the image. Change the `taskRef` to specify both `tests` and `lint`.
 - ▶ Click here for a hint.

Check Your Answer

Check that your step looks like this:

- ▶ Click here for the answer.

Apply the Changes

You are now ready to apply the changes to your pipeline.

```
oc apply -f tekton/pipeline.yaml
```

You should see the message: `pipeline.tekton.dev/cd-pipeline configured`

Commit Your Changes

Because the Cloud IDE is ephemeral, it may timeout and you will lose all of your work. It is strongly recommended that you commit your changes and push them to your remote branch so that you can pull them from GitHub if needed.

```
git commit -am 'added build task'  
git push
```

Start the Pipeline

When you start the pipeline, you need to pass in the build-image parameter, which is the name of the image to build.

This will be different for every learner that uses this lab. Here is the format:

```
image-registry.openshift-image-registry.svc:5000/$SN_ICR_NAMESPACE/accounts:latest
```

Notice the variable `$SN_ICR_NAMESPACE` in the image name. This is automatically set to point to your container namespace. There is nothing for you to do here.

Make sure that the `GITHUB_ACCOUNT` environment variable is set to the name of your GitHub account, and use the `tkn pipeline start` command to run the pipeline.

```
tkn pipeline start cd-pipeline \
-p repo-url="https://github.com/$GITHUB_ACCOUNT/devops-capstone-project.git" \
-p branch=main \
-p build-image=image-registry.openshift-image-registry.svc:5000/$SN_ICR_NAMESPACE/accounts:1 \
-w name=pipeline-workspace,claimName=pipelinerun-pvc \
-s pipeline \
--showlog
```

Check that everything ran to completion (this may take a while).

Exercise 8: Add the Deploy Task to the Pipeline

You are now at the deploy step, which is the last step in your CD pipeline. For this step, you will use the OpenShift client to deploy your Docker image to an OpenShift cluster.

Knowing that you want to deploy to OpenShift, you search Tekton Hub for "openshift" and you see there is a task called `openshift-client` that will execute OpenShift commands on your cluster. You decide to use the `openshift-client` task in your pipeline to deploy your image.

First, you need to update the `pipeline.yaml` file to reference the new `openshift-client` ClusterTask that you want to use in the deploy pipeline. Then you will need to update your `deploy/deployment.yaml` file to substitute the name of the image to deploy.

In the previous steps, you simply changed the name of the reference to the task, but since the `openshift-client` task is installed as a **ClusterTask**, you need to add the statement `kind: ClusterTask` under the name so that Tekton knows to look for a **ClusterTask** and not a regular **Task**.

Your Task

1. Start by editing the `pipeline.yaml` file to add the new deploy task. Start by copying and pasting the `build` task and changing the name to `deploy`.

► Click here for a hint.

2. Next, change the `name:` of the `workspaces:` to `manifest-dir` as specified in the documentation of the `openshift-client` on Tekton Hub.

► Click here for a hint.

3. Then, change the `taskRef` from `buildah` to `openshift-client` and leave the `kind:` as `ClusterTask` to indicate that this is a ClusterTask.

► Click here for a hint.

4. Now, change the `IMAGE` parameter to the name `SCRIPT` and specify the `oc` command to apply your Kubernetes manifests from the previous lab. Before you do, you will need to modify your `deployment.yaml` file to have an image name placeholder and use `sed` to change it to the image name for the pipeline before execution.

Here is the script to use (you will modify the `deployment.yaml` file in a later step):

```
echo "Updating manifest..."
sed -i "s|IMAGE_NAME_HERE|$(params.build-image)|g" deploy/deployment.yaml
cat deploy/deployment.yaml
echo "Deploying to OpenShift..."
oc apply -f deploy/
oc get pods -l app=accounts
```

► Click here for a hint.

5. Finally, specify that this task should `runAfter:` the `build` task.

► Click here for a hint.

Check Your Answer

If you changed everything correctly, the full deploy task in the pipeline should look like this:

► Click here for the answer.

Apply the Pipeline

Apply the same changes you just made to `pipeline.yaml` to your cluster:

```
kubectl apply -f tekton/pipeline.yaml
```

If it applied correctly, you can commit your changes. That doesn't mean it will run without errors, but it's a good idea to commit what you have before you start debugging.

Update Your Deployment Manifest

There is one last thing that you need to do before you can start the pipeline. You must edit your `deploy/deployment.yaml` file to specify a placeholder for the name of the image that you just built in the `build` step.

[Open `deployment.yaml` in IDE](#)

Find the `image:` tag and change its value to `IMAGE_NAME_HERE`.

For example:

```
spec:
  template:
    spec:
      containers:
        - image: IMAGE_NAME_HERE
          name: accounts
```

Commit Your Changes

Because the pipeline is going to clone your Git repository and you changed `deploy/deployment.yaml`, you must push all of your code changes back to your remote branch in GitHub. You will also need to use this branch name when you run the pipeline. If you followed the instructions, the current branch should be: `cd-pipeline`.

```
git commit -am 'added deploy task'
git push
```

Exercise 9: Apply Changes and Run the Pipeline

Now that you have built out the entire pipeline, it is time to run it. If you were in a persistent environment you would set up a webhook in GitHub so that whenever your code changes, the CD pipeline would run after all the tests and checks had passed.

Unfortunately, the lab environment is ephemeral and the best we can do is start the pipeline manually.

Start the Pipeline

When you start the pipeline, you will need to use the name of the branch that you are working on, which should be `cd-pipeline`.

Now, start the pipeline to see your new deploy task run. Use the Tekton CLI pipeline start command to run the pipeline, passing in the parameters `repo-url`, `branch`, and `build-image` using the `-p` option. Specify the workspace `pipeline-workspace` and persistent volume claim `pipelinerun-pvc` using the `-w` option:

Your Task

1. First, make sure that you have set the environment variable `GITHUB_ACCOUNT` to contain the name of your GitHub account.
2. Use the `tkn pipeline start` command to start the pipeline and see if your deploy task works.

```
tkn pipeline start cd-pipeline \
  -p repo-url="https://github.com/$GITHUB_ACCOUNT/devops-capstone-project.git" \
  -p branch=cd-pipeline \
  -p build-image=image-registry.openshift-image-registry.svc:5000/$SN_ICR_NAMESPACE/accounts:1 \
```

```
-w name=pipeline-workspace,claimName=pipelinerun-pvc \
-s pipeline \
--showlog
```

Check the Deployment

Now, check to see if the deployment is running. Use the `oc` command to check that your deployment is in a running state.

```
oc get all -l app=accounts
```

Evidence

For both **Option 1: AI-Graded Submission and Evaluation** and **Option 2: Peer-Graded Submission and Evaluation**, you must provide the complete logs of your Tekton pipeline run.

1. Open your terminal and run the following command to capture the full pipeline logs into a file named `pipelinerun.txt`:

```
tkn pipelinerun logs -L > pipelinerun.txt
```

2. Open the generated `pipelinerun.txt` file in the Cloud IDE editor and copy its full contents into a file with the **same name** on your local machine.
This file will be uploaded as your evidence.

Exercise 10: Make a Pull Request

Now that you have a working deployment, it's time to push the Kubernetes manifests up to GitHub and make a pull request. Then, you will merge the request and move your story to Done.

Your Task

1. Use `git status` to make sure that you have committed your changes locally in the development environment.

2. If there are uncommitted changes, commit your changes using the message `Final CD pipeline`.

3. Push your local changes to a remote branch.

Note: Use your GitHub **Personal Access Token** as your password in the Cloud IDE environment. Also, you may have to configure Git the first time you use it.

4. Make a pull request.

5. Once the test cases pass, merge your pull request.

6. Move your story to the Done column on your kanban board.

7. Pull the last code down to your development environment and delete your old branch.

```
git checkout main  
git pull  
git branch -d cd-pipeline
```

Evidence

For both **Option 1: AI-Graded Submission and Evaluation** and **Option 2: Peer-Graded Submission and Evaluation**, you must provide a screenshot showing the story marked as complete.

1. Open your Kanban board and take a screenshot showing the story “**Create a CD pipeline to automate deployment to Kubernetes**” in the **Done** column.
Save the screenshot as `cd-pipeline-done.jpeg` or `cd-pipeline-done.png`.

Conclusion

Congratulations! You have built a CD pipeline using Tekton to clone, test, lint, build, and deploy your microservice to an OpenShift/Kubernetes cluster using YAML manifests.

Next Steps

You have completed the DevOps Capstone Project. It's now time to upload all of your evidence for the entire project.

Author

Tapas Mandal
[John J. Rofrano](#)

Other Contributor(s)